

Alpenglow Assist Training Evaluation Report

Training status: blocked_pending_model_provider_decision

Dataset records: 108 synthetic examples

Dataset integrity: upload JSONL contains one object per line; shorter reviewer excerpts are previews only.

Split: train 96 / validation 12

Topic distribution: account_access=18, billing_basics=18, onboarding=18, safe_escalation=18, troubleshooting

Candidate cleaning: retained 108 / excluded 12

Excluded anomalies: contradictory_answer, direct_answer_to_escalation_case, exact_duplicate, incomplete_re

Duplicate leakage check: no duplicate user messages across train and validation.

Tone check: all retained examples use concise_support.

Safety check: escalation cases use handoff language, not direct account or billing actions.

Training iterations: not run; pending human model/provider/runtime decision.

Held-out evaluation prompts: six prompts cover onboarding, troubleshooting, account access, billing basics, wo

Held-out results: blocked because no checkpoint exists.

Failure cases: no model failures observed; expected risk areas are verbosity, missed escalation, and billing wor

Readiness verdict: blocked_pending_model_provider_decision. Upload-ready data exists; deployment is not ap